

Manipulation of noise in generative modeling

Supervisor: Luca Eyring, Akata Lab

Student: Artem Serebriakov

Introduction

Common issues of Text-to-Image models:

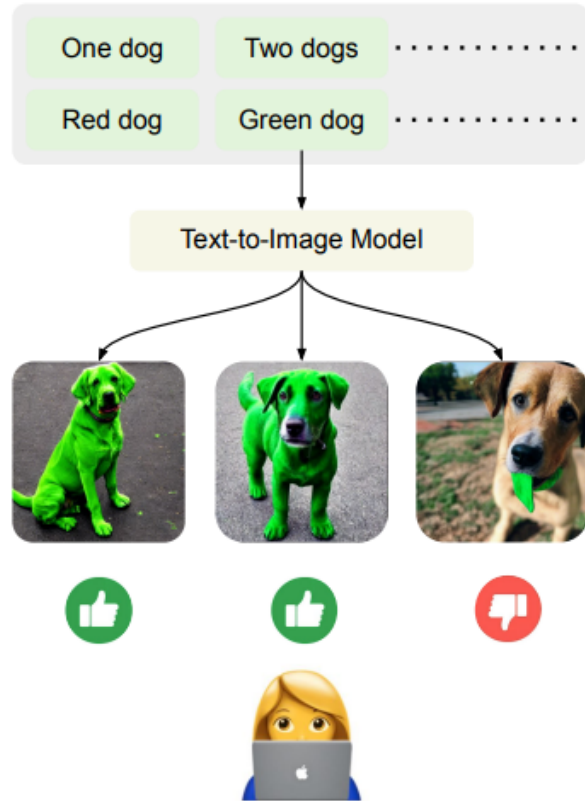
- Incorrect text rendering;
- Poor attribute binding;
- Image inaccuracies.

Approaches to address these issues:

- Employ enhanced language encoders and larger diffusion models;
- Fine-tune T2I models.

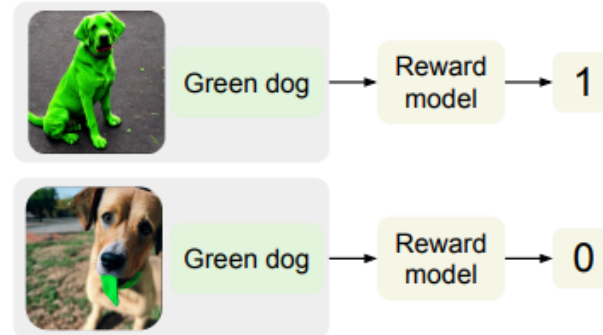
Human preference reward models

Step 1. Collecting human data

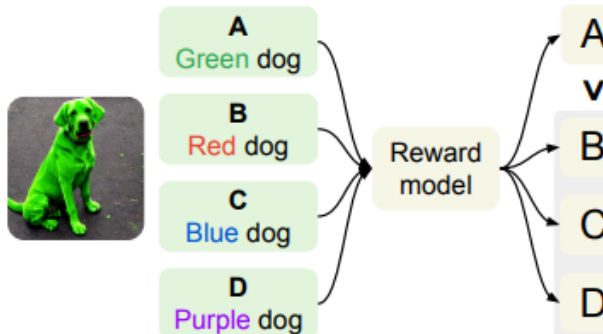


Step 2. Learning reward function

(a) Predicting human feedback



(b) Auxiliary objective: prompt classification



Step 3. Updating text-to-image model

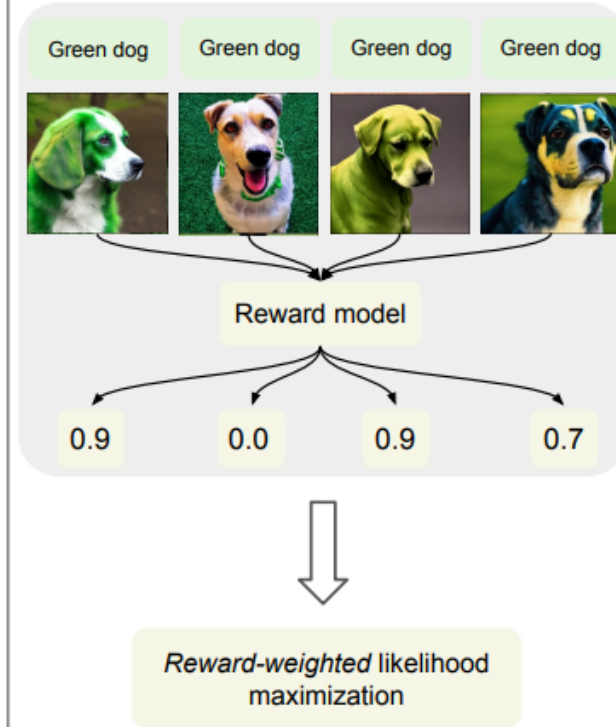


Figure 1. The steps in our fine-tuning method. (1) Multiple images sampled from the text-to-image model using the same text prompt, followed by collection of (binary) human feedback. (2) A reward function is learned from human assessments to predict image-text alignment. We also utilize an auxiliary objective called prompt classification, which identifies the original text prompt within a set of *perturbed* text prompts. (3) We update the text-to-image model via reward-weighted likelihood maximization.

Motivation for ReNO

- Can we improve T2I models at inference time without any fine-tuning?
- Does the initial noise have a huge effect on generations?
- How can we modify the initial noise to get desirable outputs?


ReNO: Enhancing One-step Text-to-Image Models through Reward-based Noise Optimization

p : "A **blue scooter** is parked **near a curb in front of a green vintage car**"

Reward Models

$$\mathcal{R}_\psi^0, \mathcal{R}_\psi^1, \dots$$

Initial
Latent Noise



One-Step
T2I Model
 $G_\theta(\epsilon, p)$



\mathcal{R}_ψ^0
ImageReward

0.35


\mathcal{R}_ψ^1
HPSv2

0.28

Iteratively update Noise to
maximize Reward Scores

$$\epsilon^{t+1} = \epsilon^t + \nabla_{\epsilon^t} \sum_i \lambda_i \mathcal{R}_\psi^i(G_\theta(\epsilon^t, p), p)$$

ReNO Updated
Latent Noise



One-Step
T2I Model
 $G_\theta(\epsilon, p)$



\mathcal{R}_ψ^0
ImageReward

1.95

\mathcal{R}_ψ^1
HPSv2

0.36

Background: One-Step Diffusion Models

$$\mathbf{x}_t = \alpha_t \mathbf{x}_0 + \sigma_t \boldsymbol{\varepsilon}, \quad (1)$$

where $G_\theta(\boldsymbol{\varepsilon}, \mathbf{p}) = \mathbf{x}_0$, Text-to-Image generative model $G_\theta(\boldsymbol{\varepsilon}, \mathbf{p})$, noise vector $\boldsymbol{\varepsilon} \sim \mathcal{N}(0, \mathbf{I})$, prompt \mathbf{p} , parameters θ , α_t is a decreasing and σ_t is an increasing function of t .

Forward SDE: $d\mathbf{x}_t = \mathbf{u}(\mathbf{x}_t, t) dt + g(t) d\mathbf{w}_t$

where $\mathbf{u}_t(x_t, t)$ — drift, \mathbf{w}_t — Wiener process, $g(t)$ — diffusion schedule.

$$\text{Reverse-time SDE: } d\mathbf{x}_t = [\mathbf{u}(\mathbf{x}_t, t) - g(t)^2 \mathbf{s}(\mathbf{x}_t, t)] dt + g(t) d\bar{\mathbf{w}}_t, \quad (2)$$

where $\mathbf{x}_T = \boldsymbol{\varepsilon} \sim \mathcal{N}(0, \mathbf{I})$, $\mathbf{x}_0 \sim p_0(\mathbf{x})$, $\mathbf{s}(\mathbf{x}, t) = \nabla \log p_t(\mathbf{x})$ is the score function.

$$\text{Denoising loss: } \mathcal{L}_s(\theta) = \mathbb{E}_{\mathbf{x}_0 \sim p(\mathbf{x}_0), \boldsymbol{\varepsilon} \sim \mathcal{N}(0, \mathbf{I}), t \sim \mathcal{U}(0, T)} [\|\sigma_t \mathbf{s}_\theta(\mathbf{x}_t, t) + \boldsymbol{\varepsilon}\|^2]. \quad (3)$$

where parameterized score $\mathbf{s}_\theta(\mathbf{x}_t, t)$.

Distillation

- Adversarial Diffusion Distillation (ADD) combines score distillation with an adversarial loss and is employed to train SD-Turbo based on SD 2.1 as a teacher and SDXL-Turbo based on SDXL;
- Diffusion Matching Distillation (DMD) additionally leverages a distributional loss based on an approximated KL divergence and is applied for PixArt- α DMD;
- Trajectory Segmented Consistency Distillation (TSCD) introduces a progressive segment-wise consistency distillation loss to train HyperSDXL [65] with reward fine-tuning.

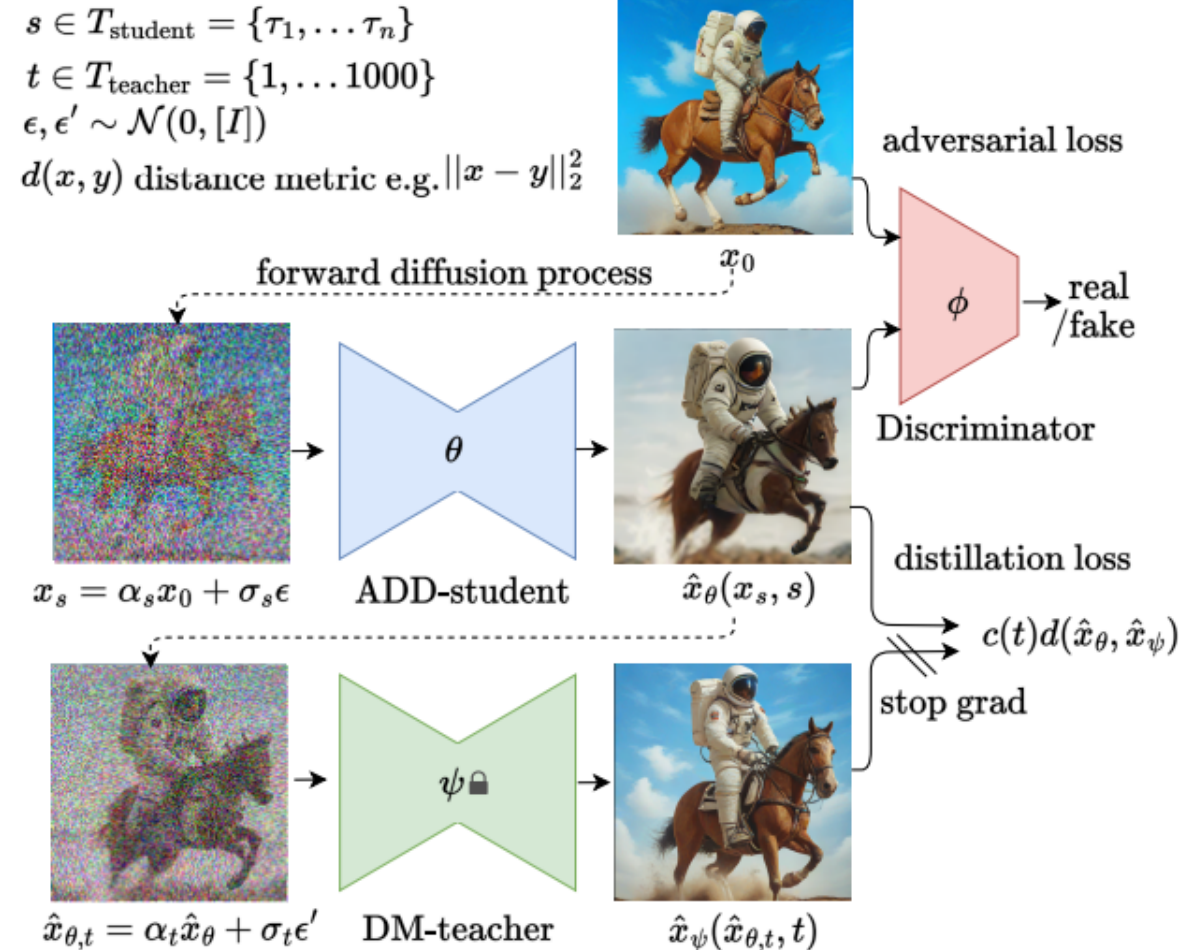


Figure 2. **Adversarial Diffusion Distillation.** The ADD-student is trained as a denoiser that receives diffused input images x_s and outputs samples $\hat{x}_\theta(x_s, s)$ and optimizes two objectives: a) adversarial loss: the model aims to fool a discriminator which is trained to distinguish the generated samples \hat{x}_θ from real images x_0 . b) distillation loss: the model is trained to match the denoised targets \hat{x}_ψ of a frozen DM teacher.

Initial Noise Optimization and Regularization

Optimization problem: $\epsilon^* = \arg \max_{\epsilon} \mathcal{C}(G_{\theta}(\epsilon, p))$.

where $\mathcal{C} : \mathbb{R}^{H \times W \times c} \rightarrow \mathbb{R}$ — criterion function, prompt p , noise ϵ .

Regularized criterion function: $\mathcal{C}(\mathbf{x}_0, \epsilon) = \tilde{\mathcal{C}}(\mathbf{x}_0) + K(\epsilon)$,

where $K(\epsilon)$ — regularization term.

$$\tilde{\mathcal{C}}(\mathbf{x}_0) = \sum_{i,j} \mathbf{x}_0^{i,j,c} - \mathbf{x}_0^{i,j,\bar{c}_1} - \mathbf{x}_0^{i,j,\bar{c}_2}, \quad (5)$$

where $\mathbf{x}_0^{i,j,c}$ — channel c of the pixel at (i, j) .

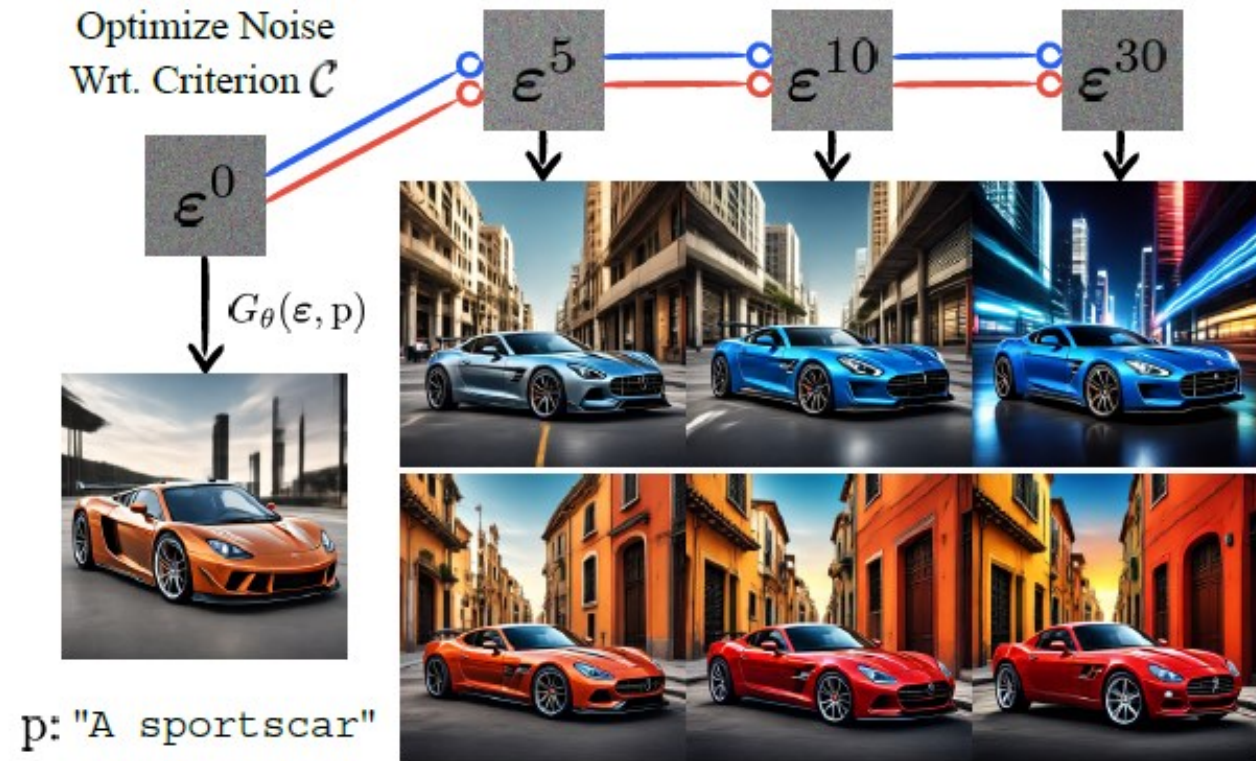


Figure 3: Initial noise optimization for one-step G_{θ} HyperSDXL with two color channel criterions (5).

ReNO Reward Criterion

Reward-based criterion function for Noise Optimization (ReNO):

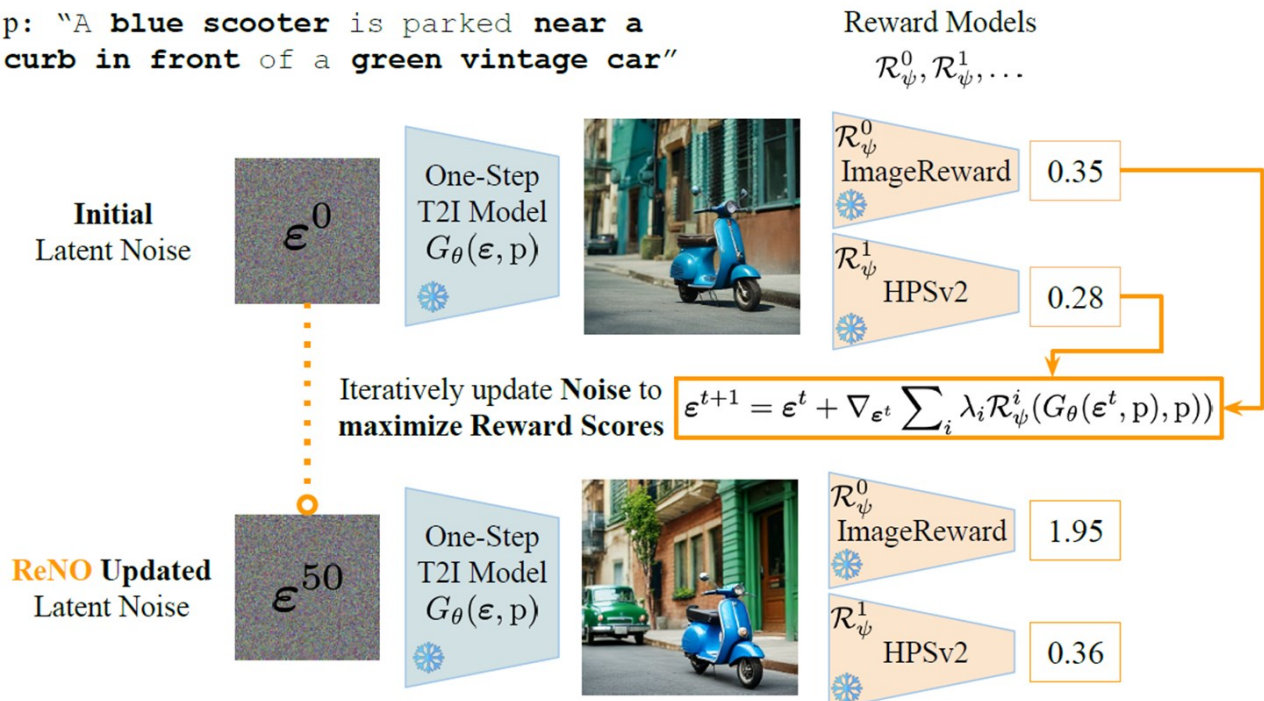
$$\tilde{\mathcal{C}}(\mathbf{x}_0, \mathbf{p}) = \sum_i^n \lambda_i \mathcal{R}_\psi^i(\mathbf{x}_0, \mathbf{p}), \quad (6)$$

where pre-trained reward models $\mathcal{R}_\psi^0, \dots, \mathcal{R}_\psi^n$, λ_i denotes the weighting for reward model \mathcal{R}_ψ^i .

$$\text{Gradient ascent: } \epsilon^{t+1} = \epsilon^t + \eta \nabla_{\epsilon^t} [K(\epsilon^t) + \sum_i^n \lambda_i \mathcal{R}_\psi^i(G_\theta(\epsilon^t, \mathbf{p}), \mathbf{p})], \quad (7)$$

where η is the learning rate.

\mathbf{p} : "A blue scooter is parked near a curb in front of a green vintage car"



Algorithm 1 ReNO

Input: \mathbf{p} (prompt), G_θ (One-Step T2I Model), $\mathcal{R}_\psi^{0,1,\dots,n}$ (Reward Functions), $\lambda_{0,1,\dots,n}$ (Reward Weights), m (# Optimization Steps), η (Learning Rate), λ_{reg} (Regularization Strength)

Initialize $v_{-1} = 0.0$, $\epsilon^0 = \mathcal{N}(0, \mathbf{I})$, $R^* = -\text{inf}$.

for $t = 0$ **to** m **do**

 Generate image $\mathbf{x}_0^t = G_\theta(\epsilon^t, \mathbf{p})$

 Compute reward-based criterion $R^t = \sum_i^n \lambda_i \mathcal{R}_\psi^i(\mathbf{x}_0^t, \mathbf{p})$

$\text{grad}_t = \nabla_{\epsilon^t} [\lambda_{\text{reg}} K(\epsilon^t) + R^t]$

$\text{grad}_t = \text{GradNormClip}(\text{grad}_t, 0.1)$

$v_t = 0.9 \cdot v_{t-1} + \eta \cdot \text{grad}_t$

$\epsilon^{t+1} = \epsilon^t + v_t$

if $R^t > R^*$ **then**

$\mathbf{x}_0^* = \mathbf{x}_0^t$, $R^* = R^t$

end if

end for

return \mathbf{x}_0^*

Experimental Setup

One-Step T2I Models			
SD-Turbo	SDXL-Turbo	PixArt- α DMD	HyperSDXL

Benchmarks		
T2I-CompBench	GenEval	Parti-Prompts

Reward Models			
PickScore	HPSv2	ImageReward	CLIPScore

Effect of Reward Models

Table 1: SD-Turbo evaluated on the attribute binding categories of T2I-CompBench and the LAION aesthetic score predictor [74] for different reward models.

Reward	Attribute Binding			Aesthetic \uparrow
	Color \uparrow	Shape \uparrow	Texture \uparrow	
SD-Turbo	0.5513	0.4448	0.5690	5.647
+ CLIPScore	0.6625	0.5501	0.6621	5.475
+ HPSv2	0.6443	0.5451	0.6859	5.752
+ ImageReward	0.7720	0.6104	0.7334	5.611
+ PickScore	0.6341	0.5069	0.6242	5.711
+ All	0.7830	0.6244	0.7466	5.704

Table 2: **Quantitative Results on T2I-CompBench.** ReNO combined with (1) PixArt- α DMD [11, 12, 91], (2) SD-Turbo [72], (3) SDXL-Turbo [72], HyperSD [65] demonstrates superior compositional generation ability in both attribute binding, object relationships, and complex compositions. The best value is bolded, and the second-best value is underlined. Multi-step results taken from [12, 20].

Model	Attribute Binding			Object Relationship		Complex \uparrow
	Color \uparrow	Shape \uparrow	Texture \uparrow	Spatial \uparrow	Non-Spatial \uparrow	
SD v1.4	0.38	0.36	0.42	0.12	0.31	0.31
SD v2.1	0.51	0.42	0.49	0.13	0.31	0.34
SDXL	0.64	0.54	0.56	0.20	0.31	0.41
PixArt- α	0.69	0.56	0.70	0.21	0.32	0.41
DALL-E 2	0.57	0.55	0.64	0.13	0.30	0.37
DALL-E 3	0.81	0.68	0.81	-	-	-
(1) PixArt- α DMD	0.38	0.34	0.47	0.19	0.30	0.36
(1) + ReNO (Ours)	0.64	0.57	0.72	0.25	0.31	0.46
(2) SD-Turbo	0.55	0.44	0.57	0.17	0.31	0.41
(2) + ReNO (Ours)	0.78	0.62	0.75	0.22	0.32	0.48
(3) SDXL-Turbo	0.61	0.44	0.60	0.24	0.31	0.43
(3) + ReNO (Ours)	0.78	0.60	0.74	0.26	0.31	0.47
(4) HyperSDXL	0.65	0.50	0.65	0.25	0.31	0.46
(4) + ReNO (Ours)	<u>0.79</u>	<u>0.63</u>	<u>0.77</u>	0.26	0.31	0.48

Quantitative Results

Table 3: **Quantitative Results on GenEval.** ReNO combined with (1) PixArt- α DMD [11, 12, 91], (2) SD-Turbo [72], (3) SDXL-Turbo [72], HyperSDXL [65] improves results across all categories. The best value is bolded, and the second-best value is underlined. Multi-step results taken from [20].

Model	Mean \uparrow	Single \uparrow	Two \uparrow	Counting \uparrow	Colors \uparrow	Position \uparrow	Color Attribution \uparrow
SD v2.1	0.50	0.98	0.51	0.44	0.85	0.07	0.17
SDXL	0.55	0.98	0.74	0.39	0.85	0.15	0.23
IF-XL	0.61	0.97	0.74	0.66	0.81	0.13	0.35
PixArt- α	0.48	0.98	0.50	0.44	0.80	0.08	0.07
DALL-E 2	0.52	0.94	0.66	0.49	0.77	0.10	0.19
DALL-E 3	<u>0.67</u>	0.96	<u>0.87</u>	0.47	0.83	0.43	0.45
SD3 (8B)	0.68	0.98	0.84	<u>0.66</u>	0.74	<u>0.40</u>	<u>0.43</u>
(1) PixArt- α DMD	0.45	0.95	0.38	0.46	0.76	0.05	0.09
(1) + ReNO (Ours)	0.59	0.98	0.72	0.58	0.85	0.15	0.27
(2) SD-Turbo	0.49	0.99	0.51	0.38	0.85	0.07	0.14
(2) + ReNO (Ours)	0.62	1.00	0.82	0.60	0.88	0.12	0.33
(3) SDXL-Turbo	0.54	1.00	0.66	0.45	0.84	0.09	0.20
(3) + ReNO (Ours)	0.65	1.00	0.84	0.68	<u>0.90</u>	0.13	0.35
(4) HyperSDXL	0.56	1.00	0.76	0.43	0.87	0.10	0.21
(4) + ReNO (Ours)	0.65	1.00	0.90	0.56	0.91	0.17	0.33

User Study Results and Computational Cost of ReNO




Figure 4: User Study Results for ReNO

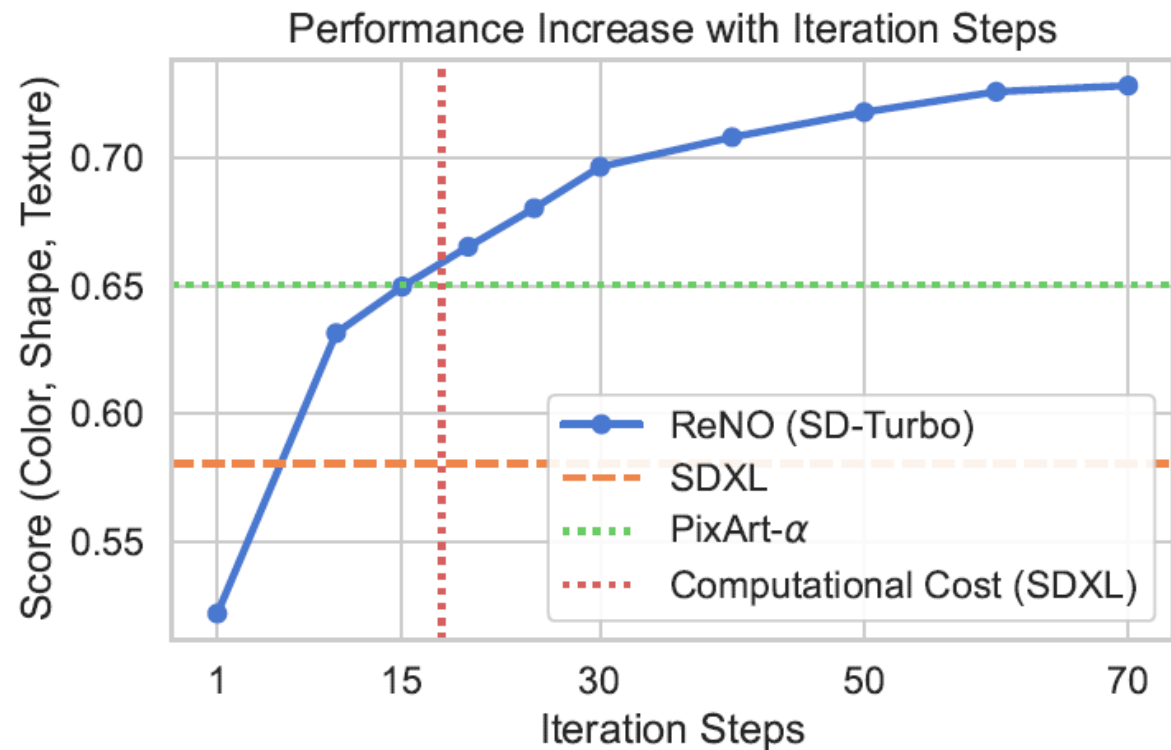


Figure 5: Attribute binding results on T2I-CompBench with varying number of iterations.

Limitations


- Despite using different image generation models of varying architectures and sizes, they converge to similar performance on both T2I-Compbench and GenEval;
- The amount of needed GPU VRAM is significantly higher when using ReNO;
- ReNO is designed for distilled diffusion models.

Further analysis

- Goal: understand how the optimization process changes the initial noise;
- Metrics: L2 distance, coherence function;
- Parti-Prompts dataset (a set of diverse 1.6k prompts).

Analysis of the distances between noises and images

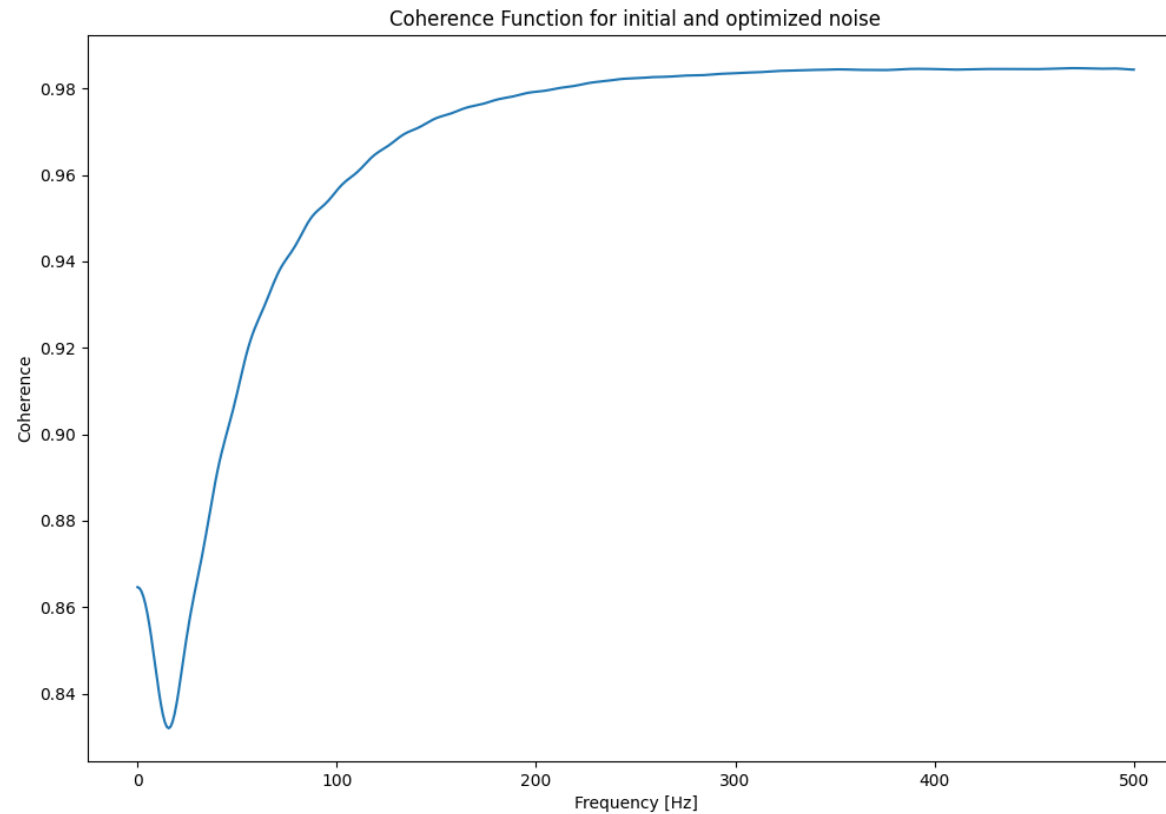
- Mean L2 distance between initial and optimized noises: 27.0122;
- Mean L2 distance between two random noises: 180.7667;
- Mean L2 distance between initial noise and generated image: 123.3377;
- Mean L2 distance between optimized noise and generated image: 142.3693;
- Mean L2 distance between random noise and random image: 159.0550.



Conclusion:

- Mean L2 distance between the initial and optimized noises is much less than distance between random noises;
- Mean L2 distance between the initial noise and image is smaller than the distance between the optimized ones.

Coherence analysis of the initial and optimized noises




- Coherence function measures the frequency-dependent correlation between two noises;
- Coherence values range from 0 (no correlation) to 1 (perfect correlation) and are computed across different frequencies;
- Low frequencies represent the "smooth" or slowly varying parts of the signal;
- High frequencies correspond to the fast-changing or detailed parts of the signal;


Conclusion: ReNO primarily modifies the slower-varying components while preserving the rapid fluctuations.

Visualization of the difference between the initial and optimized noises

Prompt: "A red dog and a green cat"



Prompt: "Oil painting of a giant robot made of sushi, holding chopsticks."



Conclusion: No visual difference between the initial and optimized noises.

Findings

- Mean L2 distance between the initial and optimized noises is much less than distance between random noises;
- Mean L2 distance between the initial noise and image is smaller than the distance between the optimized ones;
- Coherence function is lower at low frequencies and it is increasing at higher frequencies;
- No visual difference between the initial and optimized noises.

Thank you for your attention!